

Product Developer — Take-Home Task

Role: Product Developer (AI Products)

Company: Phaze AI

Time limit: 48 hours from the moment you receive this

Effort expected: 8–12 focused hours. Do not burn the full 48.

Why this task exists

We build AI products for businesses. The job is not "write code to spec" — it's **take a fuzzy business problem, decide what the product should be, and ship something a real customer could use.**

This task is deliberately under-specified. Part of what we're grading is the decisions you make when nobody tells you what to do.

You are expected to use AI tools (Claude Code, Cursor, whatever you use). That's the job. We're not testing whether you can write a for-loop from memory.

The problem

Search is moving from Google to AI. People now ask ChatGPT, Perplexity, Claude, and Google AI Overviews instead of typing keywords into a search bar.

That breaks the entire SEO industry's playbook. A business can rank #1 on Google and be **completely invisible** inside AI answers, because AI engines choose their sources by a different set of rules — and most businesses have no idea what those rules are.

This new discipline is called **GEO — Generative Engine Optimization** (also called AEO / LLMO / AI SEO).

Right now, almost no business knows whether they're visible in AI search. There is no easy way to check. That's the gap.

What to build

A GEO Auditor.

A user enters a business — at minimum a website URL — and gets back an audit report that tells them:

- 1 **How visible they currently are in AI search** (a score, and the evidence behind it)
- 2 **What is broken** (specific, with proof — not generic advice)
- 3 **What to fix, in what order** (prioritized by impact vs effort)

That's it. How you build it, what stack, what the report looks like — your call.

What to actually check — you decide

We are not giving you a checklist. Figuring out what to measure is the task.

Nobody has agreed on what a GEO audit should contain yet. The field is about two years old. So go research it: read what's being written about AI search visibility, look at what the early tools in this space do, poke at a few AI engines yourself and notice what makes them cite one source over another. Then decide what belongs in *your* auditor.

Two rules:

- **Go deep, not wide.** Three checks done properly beat twelve that tick boxes. A narrow tool that works beats a wide tool that half-works.
- **Be able to defend every check.** In your README, tell us why each one is in there and why the ones you skipped aren't. "It was easy to build" is an acceptable reason if you say so honestly.

If two candidates submit tools that check completely different things, that's expected. We're reading your reasoning as much as your code.

The report is the product

Read this section twice. **The report is what we're actually evaluating.** A beautiful scraper with an ugly report scores badly.

Requirements:

- **A score** with a visible breakdown. If you show a number, we must be able to see how it was calculated. No magic 73/100.
- **Every finding must carry evidence.** A vague "your site isn't optimized" is worthless. Name the exact page, show what you found, and show what it should be instead.
- **Prioritized fixes**, ordered by impact x effort. The business owner should know what to do Monday morning.
- **Written for a business owner, not an SEO consultant.** If a fix needs jargon, explain the jargon inline.
- **Copy-pasteable output wherever the fix allows it.** Don't describe the fix — hand it to them.

Format is your choice: web page, PDF, dashboard, whatever you can defend.

What we are NOT asking for

- Auth, signup, billing, user accounts — skip entirely
- A database, unless you actually need one

- Deployment to production (a working local run + a demo video is fine, though a live link is a nice touch)
 - Tests, CI, Docker — not graded
 - Mobile responsiveness — not graded
 - Handling every website on the internet. Make it work well for 3–5 real businesses.
-

Constraints

- **Build it however you want.** Any language, any framework, any model, any API, any tooling. Zero restrictions.
 - **Use AI tools freely.** Claude Code, Cursor, Copilot — expected, not penalized.
 - **Existing libraries are fine.** Do not hand-roll an HTML parser.
 - **If you mock something, label it** in the code and in your README. A clearly-labelled mock costs you nothing. A fake result presented as real is an instant no.
 - **Questions are welcome** — ask any time, it's a positive signal. But don't wait on an answer to start. Make an assumption, write it in the README, and move.
-

Deliverables

Submit all four:

- 1 **Code** — public GitHub repo (or a private one, and we'll send you the account to invite)
- 2 **README** covering:
 - How to run it (we should get it running in under 5 minutes)
 - What you chose to build and, more importantly, **what you chose to cut and why**
 - What's real vs mocked
 - What you'd build next with another week
- 1 **Three real audit reports** — run your tool against three actual businesses of your choosing and include the outputs. Not screenshots of a template. Real runs.
- 2 **A 3–5 minute Loom/screen recording** — you, walking through the tool and one report. No editing needed. Talk about your decisions, not just your features.

How to submit: reply on the same thread you received this task on.

Subject line: GEO Auditor – [Your Name]

If you can't finish everything, ship what works and be honest about what doesn't. A working 40% with a clear README beats a broken 90%.

How we score it

Area	Weight	What earns points
Research & what you chose to check	30%	You went into an unfamiliar field, figured out what actually drives AI visibility, and built checks that measure it. This is the biggest single factor. Picking obvious surface-level checks scores low even if they're implemented perfectly.
Report quality	25%	Would a business owner pay for this output? Is it specific, evidenced, prioritized, and readable?
Product judgment	20%	Did you cut the right things? Does this feel like it was built by someone who thought about the customer, not just the ticket?
Technical execution	15%	Does it run? Does it handle a messy real-world site without crashing? Is the code something a second person could extend?
Communication	10%	README and video. Can you explain a decision and defend a tradeoff?

What gets you an instant no

- Tool only works on the one URL you demoed
- Report is generic advice that would be identical for any website
- Findings with no evidence attached
- Mocked results presented as live data
- Checks you can't explain the reasoning behind
- No submission by the deadline with no heads-up

What gets you hired on the spot

- We run it on a business you've never seen and it produces something genuinely useful
- You measure something we hadn't thought of, and after seeing it we can't imagine the report without it
- You cut scope aggressively and can explain exactly why
- Your README changes how we think about this problem

One last thing

The businesses that need this have no idea anything is wrong. They think they're fine because they rank on Google.

The whole value of this product is the moment the owner sees the report and realises they have a problem — and then finds the fix sitting right underneath it.

Build for that moment.